

数据库系统原理

邵荃侠

课程内容

Chapter 1: Introduction

Chapter 2: Relation Model

Chapter 3: SQL

Chapter 4: Intermediate SQL

Chapter 5: Advanced SQL

Chapter 6: Database Design Using the E-R Model

Chapter 7: Relational Database Design

Chapter 8: Complex Data Types

Chapter 9: Application Development

Chapter 13: Data Storage Structures

Chapter 14: Indexing

Chapter 15: Query Processing

Chapter 16: Query Optimization

Chapter 17: Transactions

Chapter 18: Concurrency Control

Chapter 19: Recovery System

Chapter 2: Relation Model

- Relational data model
 - *relational data structure*
 - basic elements in relational data model, i.e. tables and attributes, and relationships among them
 - *integrity constraints*
 - constraints on attributes of schemas, e.g. value domain, type
 - constraints on dependencies among attributes of a schema
 - constraints on dependencies among attributes of different schemas
 - key, foreign key, etc.
 - *operations on the model*
 - Relational algebra

Chapter 2: Relation Model

Symbol (Name)	Example of Use
σ (Selection)选择	$\sigma \text{ salary} \geq 85000 \text{ (instructor)}$
	Return rows of the input relation that satisfy the predicate.
Π (Projection) 投影	$\Pi ID, salary \text{ (instructor)}$
	Output specified attributes from all rows of the input relation. Remove duplicate tuples from the output.
\times (Cartesian Product)	$\text{instructor} \times \text{department}$
	Output all pairs of rows from the two input relations whether or not they have the same value on all attributes that have the same name.
\cup (Union)并	$\Pi name \text{ (instructor)} \cup \Pi name \text{ (student)}$
	Output the union of tuples from the <i>two</i> input relations.
$-$ (Set Difference)集合差	$\Pi name \text{ (instructor)} - \Pi name \text{ (student)}$
	Output the set difference of tuples from the two input relations.
\bowtie (Natural Join) 自然连接	$\text{instructor} \bowtie \text{department}$
	Output pairs of rows from the two input relations that have the same value on all attributes that have the same name.

Exercises

employee (person_name, street, city)

works (person_name, company_name, salary)

company (company_name, city)

给出关系代数表达式来表示下列每一个查询：

- 1) 找出居住在“**Miami**”城市的所有员工姓名。
- 2) 找出工资在**100000**美元以上的所有员工姓名。
- 3) 找出居住在“**Miami**”并且工资在**100000**美元以上的所有员工姓名。

$$1) \pi_{person_name} \left(\sigma_{city='Miami'}(employee) \right)$$

$$2) \pi_{person_name} \left(\sigma_{salary > 100000}(works) \right)$$

$$3) \pi_{person_name} \left(\sigma_{city='Miami' \wedge salary > 100000} (employee \bowtie works) \right)$$

Chapter 3, 4, 5: SQL

- SQL includes DDL, DML, etc.
- DDL: Data Definition Language
 - Create, Drop, Alter (on schema)
- Objects of DDL
 - Data type
 - Table
 - Constraints
 - Index
 -

Chapter 3, 4, 5: SQL

- SQL includes DDL, DML, etc.
- DDL
 - Create, Drop, Alter

- Objects of DDL

- Data type
- Table
- Constraints
- Index
-

```
create table takes (  
    ID                varchar(5),  
    course_id        varchar(8),  
    sec_id           varchar(8),  
    semester         varchar(6),  
    year             numeric(4,0),  
    grade            varchar(2),  
    primary key (ID, course_id, sec_id, semester, year) ,  
    foreign key (ID) references student,  
    foreign key (course_id, sec_id, semester, year)  
    references section);
```

Chapter 3, 4, 5: SQL

- DML: Data Manipulation Language
 - Select, Insert, Delete, Update
- *Select*
 - Single Table
 - Join: Natural Join, Inner Join, [Left, Right, Full] Outer Join
 - Aggregate Function: Null Value
 - Subquery (难点! ! !)
 - With clause
 - Nested clause
- Relational algebra vs. Select

Chapter 3, 4, 5 — Nested Subqueries

- SQL provides a mechanism for the nesting of subqueries, to implement more complex query
- A subquery is a *select-from-group* expression that is nested within another query
- The nesting can be done in the following SQL query

```
select  $A_1, A_2, \dots, A_n$   
from  $r_1, r_2, \dots, r_m$   
where  $P$ 
```

as follows:

- A_i can be replaced by a subquery that generates a single value.
- r_i can be replaced by any valid subquery
- P can be replaced with an expression of the form:

$B <\text{operation}> (\text{subquery})$

Where B is an attribute and $<\text{operation}>$ to be defined later.

Chapter 3, 4, 5 — Nested Subqueries

- The subquery is often nested in the *where clause/having clause* , *from clause*
- Subquery in the *where clause/having* , perform tests for
 - *set membership* → *in*
 - *set comparisons* → *some/all*
 - *set cardinality* → *empty relations, not exist*
 - *duplicate tuples* → *unique*
- The *with* clause provides a way of defining a temporary relation whose definition is available only to the query in which the with clause occurs
- Subqueries in the *Select* clause
 - Scalar subquery is one which is used where a single value is expected

Use of “not exists” Clause

- Find all students who have taken all *courses* offered in the **Biology** *department*.

```
select distinct S.ID, S.name
from student as S
where not exists ( (select course_id
                    from course
                    where dept_name = 'Biology')
except
(select T.course_id
 from takes as T
 where S.ID = T.ID));
```

- First, nested query lists all courses offered in Biology
- Second, nested query lists all courses a particular student took

Note that $X - Y = \emptyset \Leftrightarrow X \subseteq Y$

Note: Cannot write this query using = **all** and its variants

Complex Queries using With Clause

- Find all departments where the total salary is greater than the average of the total salary at all departments

```
with dept_total (dept_name, value) as
    (select dept_name, sum(salary)
     from instructor
     group by dept_name),
dept_total_avg(value) as
    (select avg(value)
     from dept_total)
select dept_name
from dept_total, dept_total_avg
where dept_total.value > dept_total_avg.value;
```

Chapter 3, 4, 5: SQL

- Modification of the Database
 - Insert
 - Delete
 - Update
- Advanced SQL
 - Access from a programming language
 - JDBC, ODBC, ADO, ...
 - Embedded SQL


Chapter 6: Database Design and E-R Model

- DB/DBS/DBAS design process
 - DB design phases
requirement analysis, conceptual design, logical design, physical design
- The Entity-Relationship Model
 - *basic* E-R model
 - modeling elements: entity sets, relationship sets, attributes
 - constraints: mapping cardinality, participation constraint, keys
 - weak entity sets
 - removing redundant attributes in entity sets
 - E-R diagram


Chapter 6: Database Design and E-R Model

- *extended* E-R Features
 - OO features in E-R model, i.e specialization, generalization, attributes inheritance, constraints on generalization
 - aggregation: relationship among relationships
- Reduction to Relational Schemas
 - mapping elements in E-R model to that in relational models, i.e. conceptual schema → initial logical schema

Symbols Frequently Used in E-R Notation




entity set




relationship set




identifying
relationship set
for weak entity set




total participation
of entity set in
relationship




Symbols Frequently Used in E-R Notation




many-to-many
relationship




many-to-one
relationship




one-to-one
relationship




cardinality
limits




role indicator



ISA: generalization
or specialization



total (disjoint)
generalization



disjoint
generalization

Chapter 6: Database Design and E-R Model

- For each *entity set* and *relationship set*, there is a unique table which is assigned the name of the corresponding entity set or relationship set
 - Strong entity set vs. Weak entity set
- Each table has a number of columns; the columns generally corresponding to entities' or relationships' attributes, and have unique names
 - Composite attributes, Multivalued attributes,

Chapter 6: Database Design and E-R Model

- Reduction of a relationship set into tables are *strongly* dependent on the *mapping cardinality* constraint and *total/partial constraints* related to this relationship set
- A *many-to-many* relationship set is represented as a table with columns for the primary keys of the two participating entity sets, and any descriptive attributes of the relationship set
- Many-to-one and one-to-many relationship sets *that are total on the many-side* can be represented by adding an extra attribute to the “many” side, containing the primary key of the “one” side

Chapter 7, 8 — Schema Normalization

- Features of Good Relational Design
 - *Inserting* problem and *information redundancy*
 - *Deleting* problem
 - *Updating* problem and *information redundancy*
- Atomic Domains and First Normal Form
- Decomposition Using Functional Dependencies
- Functional Dependency Theory
- Algorithms for Functional Dependencies

函数
依赖
保持性

1NF



2NF



3NF



BCNF



4NF



5NF

消除非主属性对键的部分函数依赖

消除非主属性对键的传递函数依赖

消除主属性对键的部分和传递函数依赖

消除非平凡且非函数依赖的多值依赖

消除非平凡连接依赖

无损
连接性



Chapter 7, 8 — Schema Normalization

- A relational schema R is in **first normal form** if the domains of all attributes of R are atomic
- A relation schema R is in **BCNF** with respect to a set F of functional dependencies if for all functional dependencies in F^+ of the form $\alpha \rightarrow \beta$, where $\alpha \subseteq R$ and $\beta \subseteq R$, at least one of the following holds:
 - $\alpha \rightarrow \beta$ is trivial (i.e., $\beta \subseteq \alpha$)
 - α is a superkey for R
- A relation schema R is in **third normal form (3NF)** if for all: $\alpha \rightarrow \beta$ in F^+ , at least one of the following holds:
 - $\alpha \rightarrow \beta$ is trivial (i.e., $\beta \in \alpha$)
 - α is a superkey for R
 - Each attribute A in $\beta - \alpha$ is contained in a candidate key for R .

Chapter 7, 8 — Schema Normalization

- Functional Dependency Theory
- A functional dependency is **trivial** if it is satisfied by all instances of a relation
- **Transitive dependency** (传递函数依赖)
- **Partial dependency** (部分函数依赖)
- Armstrong's Axioms
- Closure of a Set of Functional Dependencies
- Closure of Attribute Sets
- Canonical Cover

Chapter 7, 8 — Schema Normalization

- Algorithms
 - Closure of a Set of Functional Dependencies
 - Closure of Attribute Sets
 - Canonical Cover
 - Lossless-join Decomposition
 - Dependency Preservation Decomposition
 - BCNF vs. 3NF Decomposition
 - Computing of Candidate Keys

Chapter 13: Storage and File Structure

- File organization
 - FLC vs. VLC
 - Sequential file, Heap file, Hash file, clustering
- Data-dictionary Storage
 - The **Data dictionary** (also called **system catalog**) stores **metadata**
- Data Buffer
 - Buffer-Replacement Policies

Organization of Records in Files

- **Heap** – a record can be placed anywhere in the file where there is space
- **Sequential** – store records in sequential order, based on the value of the search key of each record
- **Hashing** – a hash function computed on some attribute of each record; the result specifies in which block of the file the record should be placed
- Records of each relation may be stored in a separate file. In a **multitable clustering file organization** records of several different relations can be stored in the same file
 - Motivation: store related records on the same block to minimize I/O

Chapter 14: Indexing and Hashing

- Basic Concepts
- Ordered Indices
- B⁺-Tree Index
- Static Hashing
- Dynamic Hashing
- Comparison of Ordered Indexing and Hashing
- Index Definition in SQL
- Multiple-Key Access

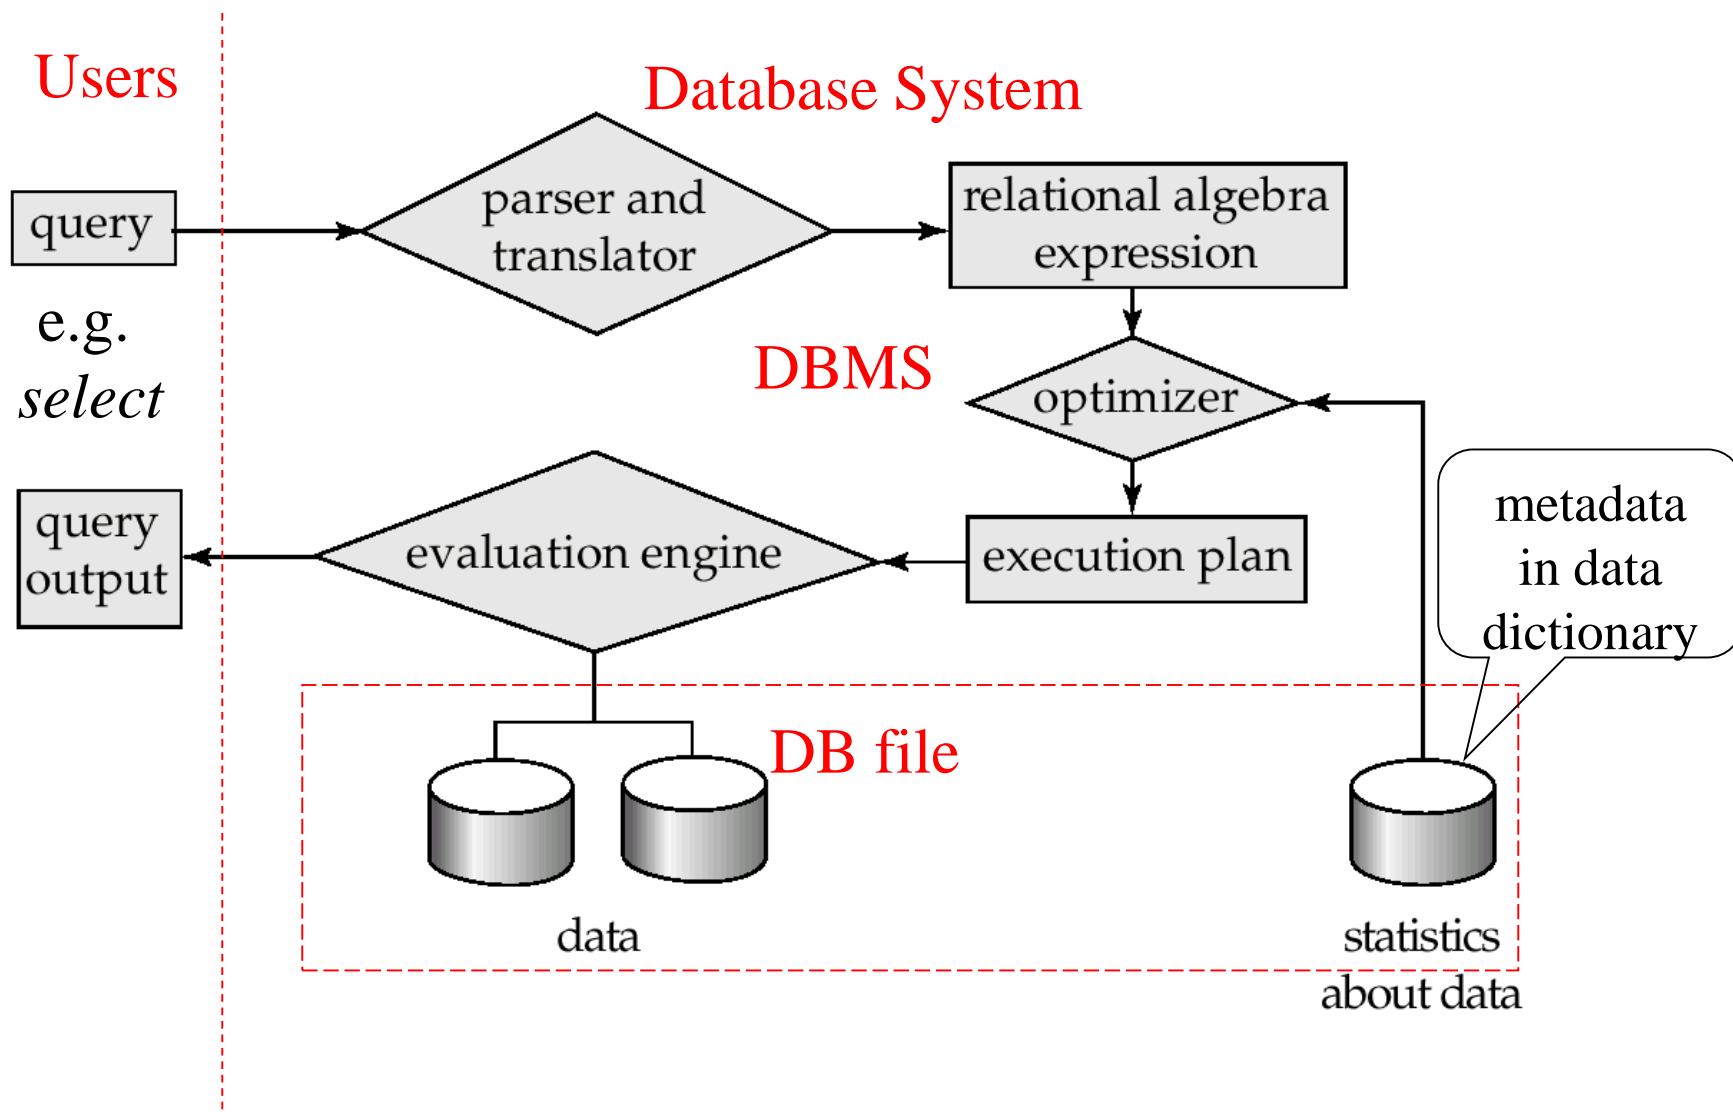
Chapter 14: Indexing and Hashing

- Basic concepts and classification of **indexing**
 - ordered indices, hash indices
- Properties/types of *ordered* indices
 - primary/clustering indices, secondary/non-clustering indices
 - dense indices, sparse indices
 - single-level indices, multi-level indices (e.g. B⁺-tree, B-tree)
- *Hash* indices
 - hash functions
 - static hash, dynamic hash

Chapter 15: Query Processing

- Basic steps in query processing
- Measures of Query Cost
 - For simplicity we just use the **number of block transfers** *from disk* and the **number of seeks** as the cost measures
- Selection Operation
 - A1-A10 file scan algorithms
- Sorting, Join Operation, Other Operations
- Evaluation of Expressions

1. Parsing and translation
2. Optimization
3. Evaluation



	Algorithm	Cost	Reason
A1	Linear Search	$t_S + b_r * t_T$	One initial seek plus b_r block transfers, where b_r denotes the number of blocks in the file.
A1	Linear Search, Equality on Key	Average case $t_S + (b_r/2) * t_T$	Since at most one record satisfies the condition, scan can be terminated as soon as the required record is found. In the worst case, b_r block transfers are still required.
A2	Clustering B ⁺ -tree Index, Equality on Key	$(h_i + 1) * (t_T + t_S)$	(Where h_i denotes the height of the index.) Index lookup traverses the height of the tree plus one I/O to fetch the record; each of these I/O operations requires a seek and a block transfer.
A3	Clustering B ⁺ -tree Index, Equality on Non-key	$h_i * (t_T + t_S) + t_S + b * t_T$	One seek for each level of the tree, one seek for the first block. Here b is the number of blocks containing records with the specified search key, all of which are read. These blocks are leaf blocks assumed to be stored sequentially (since it is a clustering index) and don't require additional seeks.
A4	Secondary B ⁺ -tree Index, Equality on Key	$(h_i + 1) * (t_T + t_S)$	This case is similar to clustering index.
A4	Secondary B ⁺ -tree Index, Equality on Non-key	$(h_i + n) * (t_T + t_S)$	(Where n is the number of records fetched.) Here, cost of index traversal is the same as for A3, but each record may be on a different block, requiring a seek per record. Cost is potentially very high if n is large.
A5	Clustering B ⁺ -tree Index, Comparison	$h_i * (t_T + t_S) + t_S + b * t_T$	Identical to the case of A3, equality on non-key.
A6	Secondary B ⁺ -tree Index, Comparison	$(h_i + n) * (t_T + t_S)$	Identical to the case of A4, equality on non-key.

Figure 15.3 Cost estimates for selection algorithms.

Chapter 16: Query Optimization

- Transformation of relational expressions
 - Rule 1- Rule 11: equivalence rules
- Choice of evaluation plans — Query optimization
 - cost-based optimization
 - *heuristic optimization*

Example Three

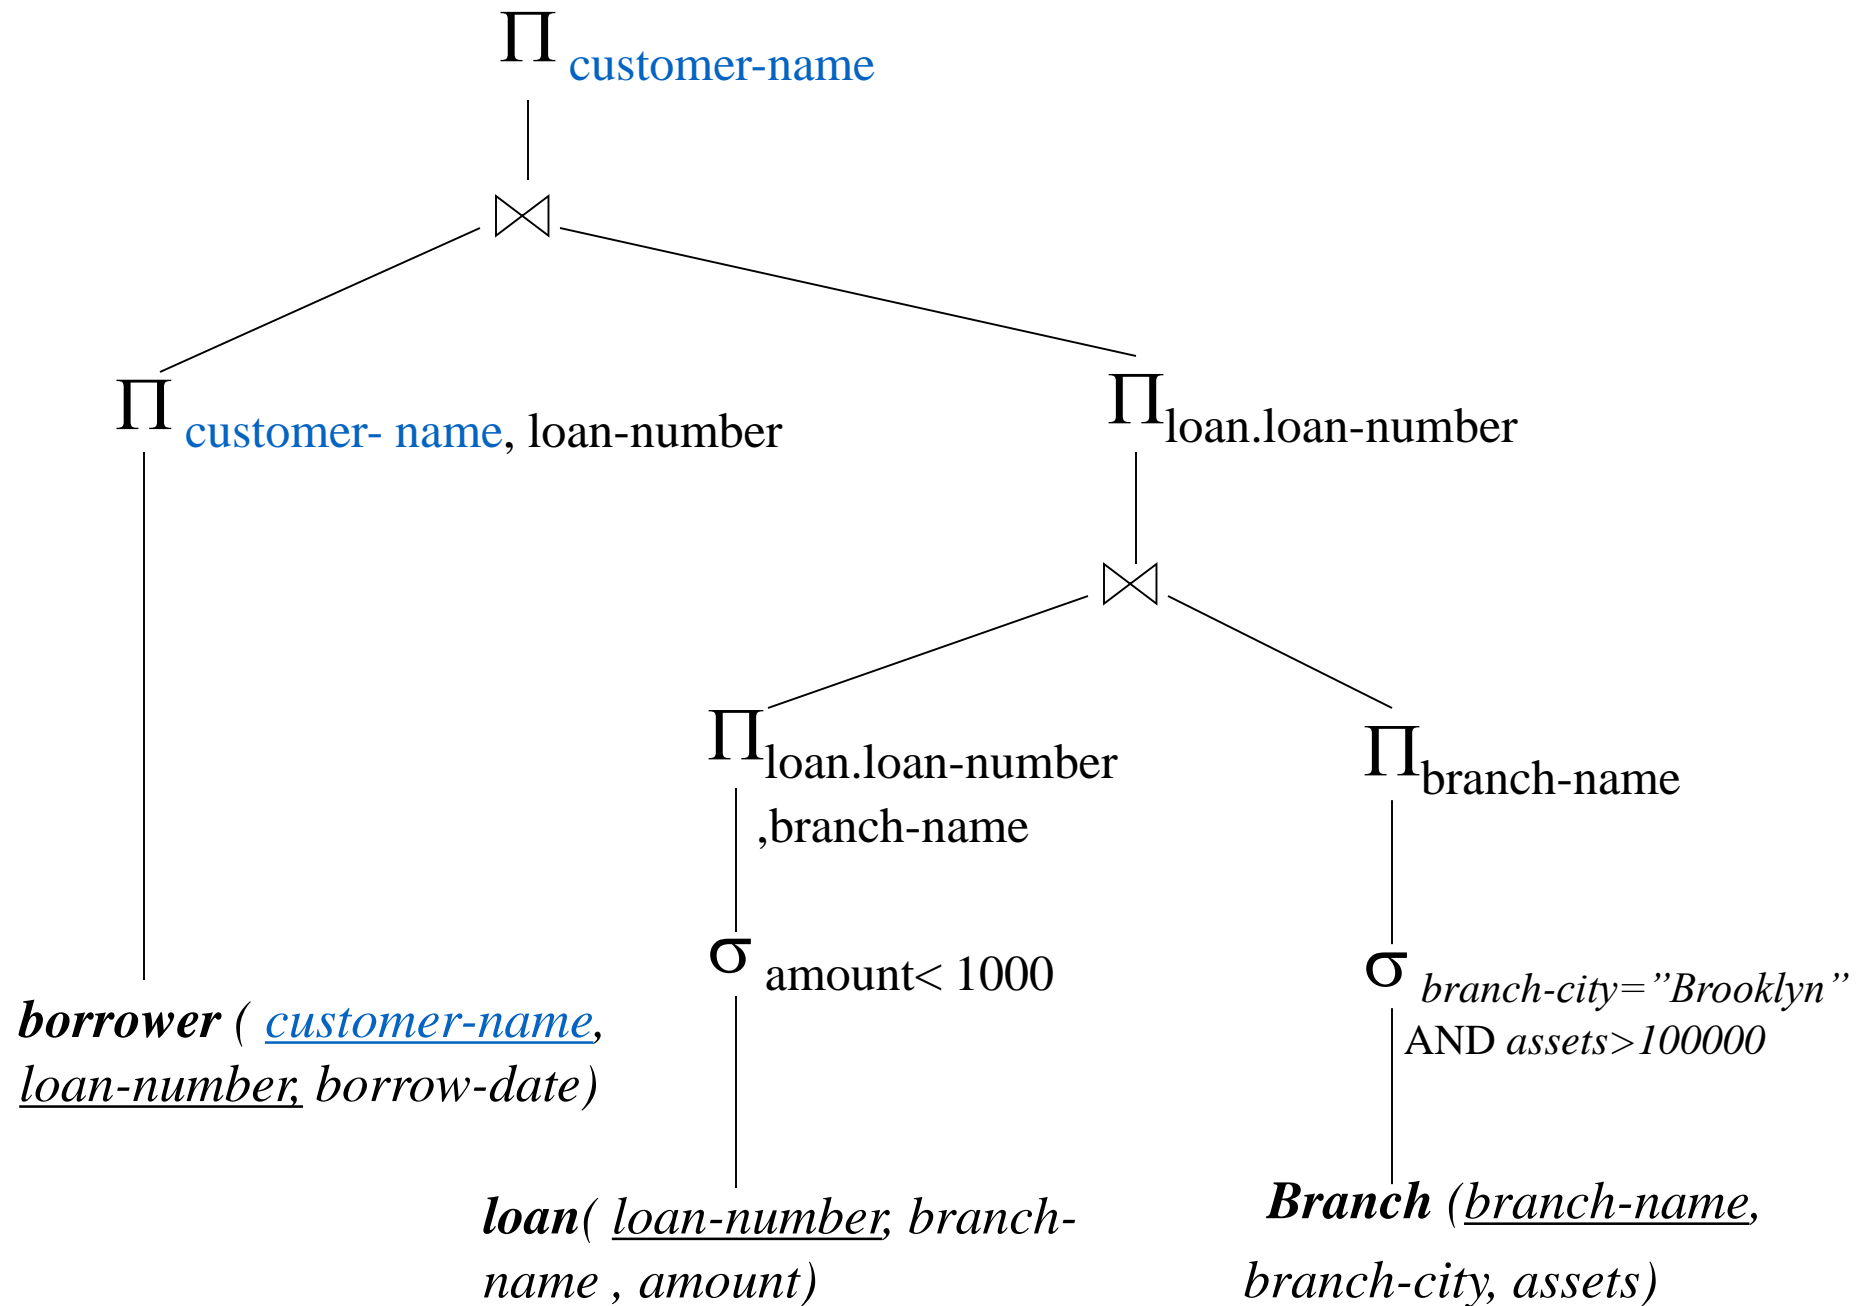
- Consider the following relations in banking enterprise database, where the primary keys are underlined
 - *branch* (*branch-name*, *branch-city*, *assets*),
 - *loan* (*loan-number*, *branch-name* , *amount*)
 - *borrower*(*customer-name*, *loan-number*, *borrow-date*)
 - *customer* (*customer-name*, *customer-street*, *customer-city*)
 - *account* (*account-number*, *branch-name*, *balance*)
 - *depositor* (*customer-name*, *account-number* , *deposit-date*)

Example Three (cont.)

- For the query “ Find the *names* of all *customers* who have an *loan* at any *branch* that is located in *Brooklyn* and have *assets* more than \$100,000, requiring that *loan-amount* is less than \$1000”
 - give an SQL statement for this query
 - given a initial query tree for the query, and convert it into an optimized query tree by means of heuristic optimization

Example Three (cont.)

- SQL
select *customer-name*
from *borrower, loan, branch*
where *loan.loan-number=borrower.loan-number*
and *branch.branch-name=loan.branch-name*
and *branch-city="Brooklyn"*
and *assets>100000 and amount<1000*



Chapter 17: Transaction

- The concept of *transaction*
- The properties of *transaction* -- **ACID**
 - 原子性 (atomicity)
 - 一致性 (consistency)
 - 独立性/隔离性 (isolation)
 - 永久性/ 持续性/操作结果永久保持性 (durability)
- The relationship between ACID and the DBS component
 - Concurrency control
 - Recovery management

Chapter 17: Transaction

- What is schedule of transaction?
- Serializability
 - Serializable schedule vs. serial schedule
 - Conflict serializability
 - Recoverable serializability
 - Cascadeless serializability
- Conflict equivalence and conflict serializability
 - How to test? → precedence graph
 - How to construct? → Topology sort

Chapter 18: Concurrency Control

- Lock-based Protocol
 - X-Lock, S-Lock
- 2 Phase Locking (2PL)
 - How to construct?
 - Strict 2PL → cascadeless and serializable schedules
 - Rigorous 2PL

复习建议

- 掌握基本概念
- 理解各类例子
- 运用算法进行推演和计算
- 复习作业题

**预祝各位同学
期末顺利！**

2021-12-27 19:30 教二401